

Lock Guide

What is the purpose of the spinlock used by a lock?

- The lock structure has a field which indicates if the lock is available or not.

```
struct Lock
{
    // ... other lock fields ...
    bool held; // Note: there are other solutions
    // ... other lock fields ...
};

void LockInit( Lock * lock )
{
    // ... lock initialization code ...
    lock->held = false; // no one owns the lock yet
    // ... other lock initialization code ...
}
```

- To take the lock, a thread must test-and-set the field which holds the lock availability

```
Acquire( Lock * lock )
{
    // ... lock code here ...
    while ( lock->held ) { ... } // this is the TEST
    lock->held = true; // this is the SET
    // ... lock code here ...
}
```

- As written above, we would have a race condition because the test-and-set operation requires mutual exclusion.
- We can use a spinlock, instead of assembly, to protect the critical section.

```
Acquire( Lock * lock )
{
    // ... non-critical section lock code here ...
    Acquire( lock->spinlock )
    // ... critical section lock code here ...
    while ( lock->held ) { ... } // this is the TEST
    lock->held = true; // this is the SET
    // ... critical section lock code here ...
    Release( lock->spinlock )
}
```

- The spinlock is only held for a brief amount of time --- only for the test-and-set.
- **DO NOT** block a thread (put it to sleep) while it owns the spinlock!

Why a loop?

- Suppose there was no test-and-set loop.

```

Acquire( Lock * lock )
{
    // ... non-critical section lock code here ...
    Acquire( lock->spinlock )
    // ... critical section lock code here ...
    if ( lock->held ) { ... block ... } // this is the TEST
    lock->held = true; // this is the SET
    // ... critical section lock code here ...
    Release( lock->spinlock )
}

```

- Given three threads, A, B and C and a lock L, which is initially available.

A	B	C
Acquire(L) // succeeds, A owns L		
Context Switch		
	Acquire(L) if (lock->held) // blocks, L is unavailable	
Context Switch		
Release(L) // L is available // B is unblocked // B placed in ready queue		
Context Switch		
		Acquire(L) // succeeds, C owns L
Context Switch		
	Exits lock->held branch lock->held = true // succeeds // B ALSO owns L!!	

- When B is unblocked and continues execution there is a possibility that the lock was taken by another thread that was scheduled to run before B. Before B can take the lock

after being unblocked it needs to ensure that lock is still available. Hence, the test is in a loop.

```
Acquire( Lock * lock )
{
    // ... non-critical section lock code here ...
    Acquire( lock->spinlock )
    // ... critical section lock code here ...
    while ( lock->held ) { ... block ... } // this is the TEST
    lock->held = true; // this is the SET
    // ... critical section lock code here ...
    Release( lock->spinlock )
}
```

Why must you release the spinlock prior to calling wchan_sleep?

- A thread that calls wchan_sleep will block.
- If that thread owns the spinlock when it blocks, then no other threads may acquire that spinlock. Those threads will end up spinning --- instead of blocking.

```
Acquire( Lock * lock )
{
    // ... non-critical section lock code here ...
    Acquire( lock->spinlock ) // other threads attempting to Acquire spin here
    // until lock blocked thread releases spinlock!
    // ... critical section lock code here ...
    while ( lock->held )
    {
        ... code here ...
        wchan_sleep( lock->wchan ) // spinlock not released before calling sleep
        ... code here ...
    }
    lock->held = true;
    // ... critical section lock code here ...
    Release( lock->spinlock )
}
```

- Must release spinlock prior to blocking to ensure that threads attempting to acquire the Lock do not spin on the spinlock protecting Acquires critical section

```
Acquire( Lock * lock )
{
    // ... non-critical section lock code here ...
    Acquire( lock->spinlock ) // threads will only spin waiting for critical section
    // and never for a blocked thread to release spinlock
    // ... critical section lock code here ...
    while ( lock->held )
    {
```

```

    ... code here ...
    Release( lock->spinlock )
    wchan_sleep( lock->wchan ) // spinlock not owned by blocked thread
    ... code here ...
}
lock->held = true;
// ... critical section lock code here ...
Release( lock->spinlock )
}

```

Why must you lock the wait channel prior to releasing the spinlock?

- Suppose we locked the wait channel after releasing the spinlock.

```

Acquire( Lock * lock )
{
    // ... non-critical section lock code here ...
    Acquire( lock->spinlock )
    // ... critical section lock code here ...
    while ( lock->held )
    {
        ... code here ...
        Release( lock->spinlock ) // release spinlock
        wchan_lock( lock->wchan ) // lock wait channel
        wchan_sleep( lock->wchan ) // block
        ... code here ...
    }
    lock->held = true;
    // ... critical section lock code here ...
    Release( lock->spinlock )
}

```

- Two threads, A and B and one lock L.

A	B
Acquire(L) // succeeds, A owns L	
Context Switch	
	Acquire(L) // ... while(L->held) { // ... Release(L->spinlock) }

Context Switch	
Release(L) // L is available	
Context Switch	
	wchan_lock(L->wchan) wchan_sleep(L->wchan) // B is asleep on L's wait channel // BUT L is AVAILABLE // Who wakes B?

- You must acquire the wait channel lock prior to releasing the spinlock to prevent this.

Exercises

1. What happens if a thread tries to acquire a lock it already owns?
2. What happens if a thread tries to release a lock it does not own?
3. Suppose N threads are blocked waiting to acquire lock L. Suppose that Release calls wchan_wakeall. What happens to the N blocked threads when Release is called?
 - a. Which of the N threads acquires L?
 - b. What happens to threads that do not acquire L?